

CS 161 Intro. To Artificial Intelligence

Week 4, Discussion 1C

Li Cheng Lan



Today's Topics

- Constraint Satisfaction Problem (CSP)
 - Formulation of CSPs
 - Backtrack Search
 - Techniques for improving CSP solution
 - Tree-structured CSPs
- Game Playing
 - Formulation as Search
 - Minimax Algorithm
 - Alpha-beta Pruning
 - Expect Minimax for Nondeterministic Games
- Propositional Logic

Today's Topics

- **Constraint Satisfaction Problem (CSP)**

- Formulation of CSPs
- Backtrack Search
- Techniques for improving CSP solution
- Tree-structured CSPs

- **Game Playing**

- Formulation as Search
- Minimax Algorithm
- Alpha-beta Pruning
- Expect Minimax for Nondeterministic Games

- **Propositional Logic**

CSP - Formulation

Components of Constraint Satisfaction Problem (CSP):

X is a set of variables, $\{X_1, \dots, X_n\}$.

D is a set of domains, $\{D_1, \dots, D_n\}$, one for each variable.

C is a set of constraints that specify allowable combinations of values.

- Each domain D_i consists of a set of allowable values $\{v_1, \dots, v_k\}$ for the corresponding X_i
- A **state** in CSP: an assignment of values to some or all variables
 - Partial assignment: assign values to only some of the variables
 - Complete assignment: every variable is assigned (otherwise partial assignment)
 - Consistent/Legal assignment: an assignment that does not violate any constraints
- A **solution** in CSP: a consistent, complete assignment

CSP – Formulation Example


Map Coloring:



- **Variables:** WA, NT, Q, NSW, V, SA, T
- **Domains:** $D_i = \{\text{red, green, blue}\}$
- **Constraints:** adjacent regions must have different colors
 - E.g. $WA \neq NT$ (if the language allows this),
 - E.g. $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), (\text{green, red}), (\text{green, blue}), \dots\}$

CSP – Formulation Example

Map Coloring:




- Solutions are assignments satisfying all constraints,
 - e.g., {WA=red, NT =green, Q=red, NSW =green, V =red, SA=blue, T =green}

Varieties in CSP

- Unary constraints: involve a single variable,
 - e.g., SA ~~≠~~ green
- **Binary** constraints: involve pairs of variables, *← look at in this class*
 - e.g., SA \neq WA
- Higher-order(Global) constraints: involve 3 or more variables,
 - e.g., Alldif (all of the variables involved in the constraint must have different values)
- Preferences (soft constraints):
 - e.g., red is better than green
 - often representable by a cost for each variable assignment
 - constrained optimization problems

Constraint Graph

- For **Binary CSP** -- each constraint relates at most two variables
- Constraint graph: nodes are variables, arcs show constraints
 - General-purpose CSP algorithms use the graph structure to speed up search.



Backtrack Search

Backtracking search: the basic uninformed algorithm for CSPs


- It's basically a depth-first search for CSPs with single-variable assignments:
 - Chooses values for one variable at a time and **backtracks** when a variable has no legal values left to assign

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

Backtrack Search - Example

- When to backtrack? → when a variable has no legal values left to assign




Backtrack Search - Example

- 4-Queens Puzzle (assume each queen in each column) as a CSP:

- Variables: Q_1, Q_2, Q_3, Q_4 -- row indices of each queen
- Domains $D_i = \{1, 2, 3, 4\}$
- Constraints:
 - $Q_i \neq Q_j$ (cannot be in same row)
 - $|Q_i - Q_j| \neq |i - j|$ (or same diagonal)

- Backtracking search:

- (Q_1, Q_2, Q_3, Q_4) :
- $(1, X, X, X) \rightarrow (1, 3, X, X) \rightarrow$ No legal assign for Q_3 , backtracking
- $(1, 4, X, X) \rightarrow (1, 4, 2, X) \rightarrow$ No legal assign for Q_4 , backtracking
- $(1, 4, 3, X) \rightarrow$ Not a legal assign for Q_3 , backtracking
- $(2, X, X, X) \rightarrow (2, 4, X, X) \rightarrow (2, 4, 1, X) \rightarrow (2, 4, 1, 3)$, Bingo! ← solution




Improving Backtracking Efficiency

3 techniques (heuristics) to improve efficiency:

- How to select unassigned **variable**?
 - Most constrained variable / Minimum remaining values (MRV)
 - Most constraining variable / Degree heuristic
- In what order should we assign **values** to each variable?
 - Least constraining value

Improving Backtracking Efficiency

- How to select unassigned **variable**?
 - Most constrained variable / Minimum remaining values (MRV):
 - choose the **variable** with the **fewest legal values**
 - If no legal values left, fail immediately



choose this
based on degree heuristic


Improving Backtracking Efficiency

- How to select unassigned **variable**?
 - Most constrained variable / Minimum remaining values (MRV)
 - Most constraining variable / Degree heuristic:
 - Choose the **variable** with the most constraints on remaining variables
 - Attempt to reduce branching factor on future choice
 - **Useful as a tie-breaker**



Improving Backtracking Efficiency

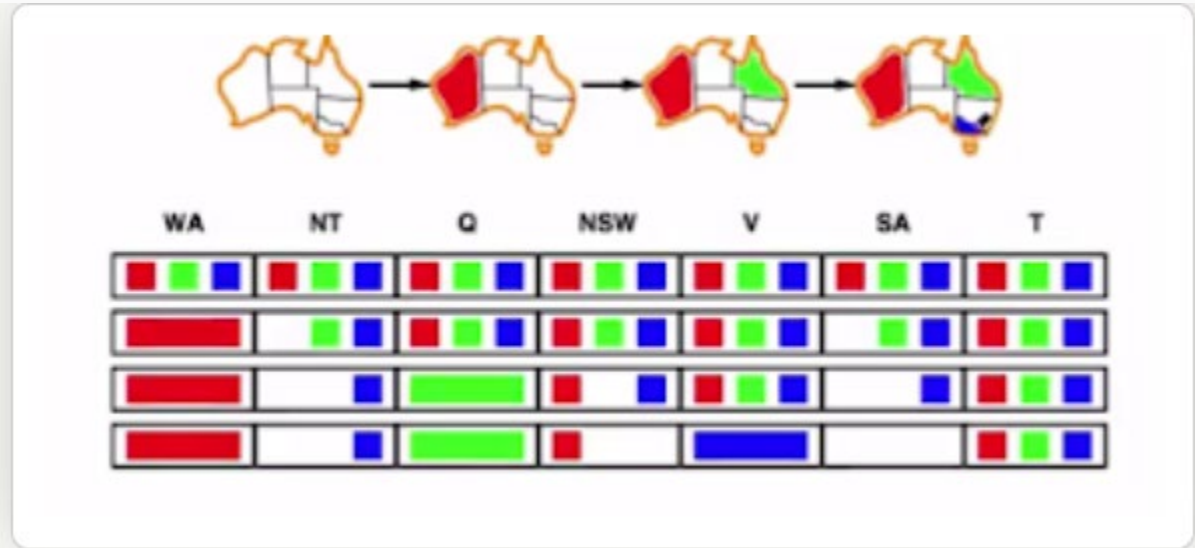
- In what order should we assign values to each variable?
 - Least constraining value:
 - Choose the **value** that **rules out the fewest values in the remaining variables**
 - Leave the maximum flexibility for subsequent variable assignments



Early Failure Detection

Two methods to detect failures by doing domain reductions:

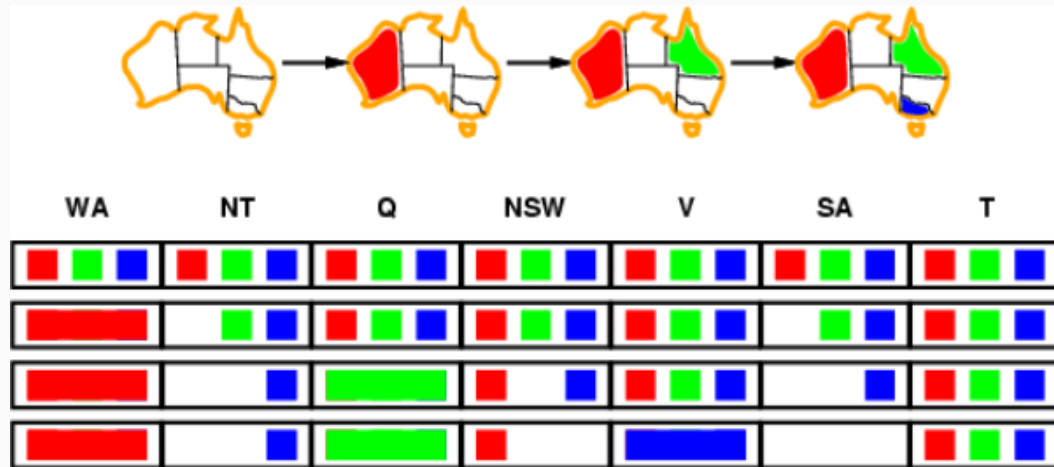
- **Forward checking**
- Arc consistency



Early Failure Detection

Two methods to detect failures by doing domain reductions:

- Forward checking
- **Arc consistency**



Arc Consistency (AC)

Evaluation of AC:

● Notations:


- n : # of variables
- d : largest domain size
- c : # of constraints

● Time complexity: $O(n^2 d^3)$

- Checking consistency of one arc: $O(d^2)$
 - Look at different combinations of values
- At most $(n^2 - n)$ arcs: $O(n^2)$ or $O(c) \leftarrow 2 * c$ constraints
- Each arc (X_i, X_j) can be inserted at most d times
 - X_i has at most d values to delete

because unidirectional/
 n variables: $\frac{2(n-1)}{2}$ arcs

$$\begin{array}{c} 1 \quad 2 \quad \dots \quad n \\ \hline 2 \times \quad 0 \quad 1 \quad \dots \quad n-1 \\ (n-1)n \end{array}$$



Arc Consistency (AC) – AC3 algorithm

Before Backtracking Search.

AC3 (not covered in class) – reduce domain size based on arc consistency before search start:

- Maintains a queue (set) of all arcs
- Pop an arbitrary arc ($X_i \leftarrow X_j$) and check D_i (the domain of X_i)
 - D_i unchanged
 - Move to next
 - D_i becomes smaller
 - Add to queue all arcs ($X_k \leftarrow X_i$) where X_k is a neighbor of X_i
 - The change in D_i might enable further reductions in the domains of D_k , even if we have previously considered X_k and D_k
 - D_i is empty
 - Fail!


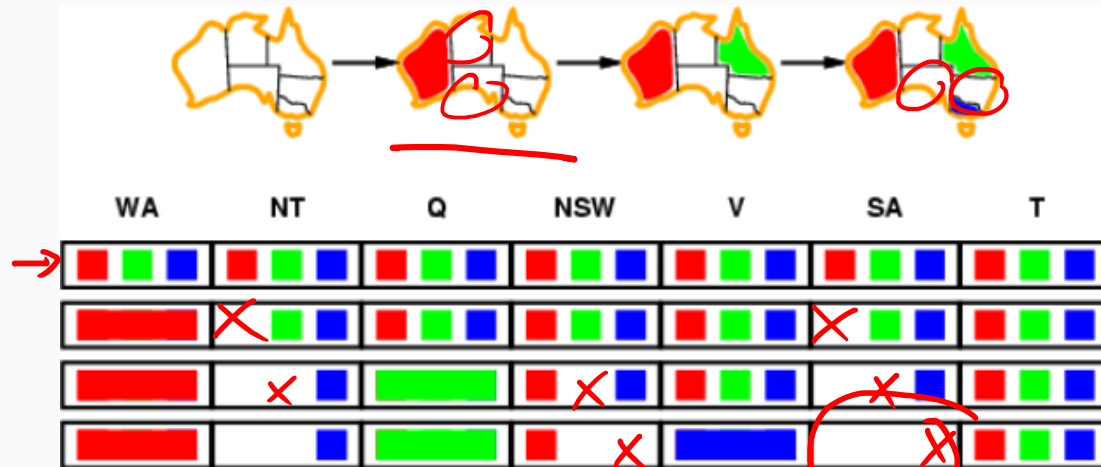
Time complexity: $O(n^2 d^3)$

Forward Checking (FC)

Inside backtracking search:

Forward Checking (during search):

- Keep track of remaining legal values for unassigned variables that are connected to current variable.
→ Variable-level arc consistency
- Terminates when any variable has no legal values
 - Then backtrack!



Forward Checking (FC)

- Doing FC every time we assign a value to a variable:
- If a variable has no legal value, do backtrack

BT: DFS-based search algorithm

*FC: technique to rule out some invalid values
→ increase efficiency of BT*

```
function BACKTRACKING-SEARCH(csp) returns solution/failure  
  return RECURSIVE-BACKTRACKING({ }, csp)
```

```
function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
```

```
  if assignment is complete then return assignment
```

```
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
```

```
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
```

```
    if value is consistent with assignment given CONSTRAINTS[csp] then
```

```
      add {var = value} to assignment
```

```
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
```

```
      if result ≠ failure then return result
```

```
      remove {var = value} from assignment
```

```
  return failure
```

Forward
checking
after this


Comparison of AC and FC

- Techniques*
- AC3: *→ check all the arcs*
- **Before** search
 - **initialization**: push **all** arcs in the queue *→ AC3*
 - **maintain the queue**: when some variable X_i 's domain size change, push $(X_k \leftarrow X_i)$ into the queue, where X_k is neighbor of X_i . *→ Do both POP and PUSH!*
- FC: *→ check arcs related to current variable.*
- **During** search
 - **initialization**: when assign value to X , push all X 's neighbor into the queue. $(neighbor \leftarrow X)$
 - **maintain the queue**: queue won't add anything after initialized. *→ ONLY POP, NO PUSH!*


Tree-structured CSPs

Theorem: If the constraint graph has no loops (tree-structured), the CSP can be solved in $O(nd^2)$ time

- n : # of variables/nodes
- d : largest domain size



Tree-structured CSPs



Algorithm with time complexity $O(nd^2)$:

n nodes $\rightarrow (n-1)$ edges $\rightarrow O(n)$ for edges

1. Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering

2. For j from n down to 2, apply $\text{Remove-Inconsistent-Values}(\text{Parent}(X_j), X_j)$ as follows:

```
function REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) returns true iff succeeds
    removed  $\leftarrow$  false
    for each  $x$  in DOMAIN[ $X_i$ ] do
        if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x,y)$  to satisfy the constraint  $X_i \leftrightarrow X_j$ 
            then delete  $x$  from DOMAIN[ $X_i$ ]; removed  $\leftarrow$  true
    return removed
```

$O(d^2)$ for each edge
 \rightarrow check are consistency

3. For j from 1 to n , assign X_j consistently with $\text{Parent}(X_j)$

Today's Topics

- Constraint Satisfaction Problem (CSP)
 - Formulation of CSPs
 - Backtrack Search
 - Techniques for improving CSP solution
 - Tree-structured CSPs
- **Game Playing**
 - Formulation as Search
 - Minimax Algorithm
 - Alpha-beta Pruning
 - Expected Minimax for Nondeterministic Games
- Propositional Logic

Types of Games


	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon monopoly
imperfect information	battleships, blind tictactoe	bridge, poker, scrabble nuclear war



Go: Perfect and Deterministic



Monopoly: Perfect, Chance Introduced



Battleship: Imperfect and Deterministic



Bridge: Imperfect, Chance Introduced

Game as a Search Problem

Can we use search strategies to win games?

- Require to make some decision when calculating the optimal decision is infeasible
- The **solution** will be a **strategy** that specifies a move for every possible opponent reply
- Challenges:
 - Very, very large search space
 - Time limits

Game as a Search Problem

- ▶ **S0**: The initial state, which specifies how the game is set up at the start.
- ▶ **PLAYER(s)**: Defines which player has the move in a state.
- ▶ **ACTIONS(s)**: Returns the set of legal moves in a state.
- ▶ **RESULT(s, a)**: The transition model, which defines the result of a move.
- ▶ **TERMINAL-TEST(s)**: A terminal test, which is true when the game is over and false otherwise.
 - ▶ States where the game has ended are called terminal states.
- ▶ **UTILITY(s, p)**: A utility function that defines the final numeric value for a game that ends in terminal state s for a player p .
 - ▶ Also called an objective function or payoff function.
 - ▶ In chess, the outcome is a win, loss, or draw, with values $+1$, 0 , or $1/2$.

↖ Utility values

Optimal Decisions in Games

How to find the optimal decision in a deterministic, perfect-information game?

Idea: choose the move with **highest achievable payoff** against the **best play of the other player**

Partial Game Tree:

- Top node is the initial state
- Giving alternating moves by MAX and MIN

Tic-tac-toe Game Tree

- Two Players: MAX: X; MIN: O

MAX (X)

MIN (O)

MAX (X)

MIN (O)


TERMINAL

Utility

-1

0

+1




Minimax Algorithm


- Imagine we are MAX
- We refer to the payoff as MINIMAX value, at each step
 - MAX wants MINIMAX value to be as big as possible
 - MIN wants MINIMAX value to be as small as possible

$\text{MINIMAX}(s) =$

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Minimax Algorithm

 : takes max value among child nodes


 : takes min value among child nodes

MAX

MIN


Order to
assign values

terminals:



Minimax Algorithm

Evaluation:

- Complete (if tree is finite)
 - Optimal (since we are against an optimal opponent)
 - Time complexity: $O(b^m)$
 - b : max # of children nodes for one parent node
 - m : max depth of the state space
 - Space complexity: $O(bm)$ → depth-first exploration
- 

Actually don't need to explore every path and compute MINIMAX for every node!

- Increase the efficiency
- Use Alpha-beta pruning

Alpha-beta Pruning

Minimax: a way of finding an optimal move in a two player game.

Alpha-beta pruning: finding the optimal minimax solution while avoiding searching subtrees of moves which won't be selected.

- α : maximum lower bound of possible solutions
- β : minimum upper bound of possible solutions
- If N is estimated value of the node, then $\alpha \leq N \leq \beta$

Alpha-beta Pruning

α : maximum lower bound of possible solutions

β : minimum upper bound of possible solutions

If N is estimated value of the node, then $\alpha \leq N \leq \beta$


Steps:

- During the search, each node carries an upper bound α a lower bound β
- Pushing bound upward:
 - When a child returns, it pushes its value onto the parent (always tighten the bound)
 - Max player will modify its lower bound, and min player will modify its upper bound
- Pushing bound downward (**both lower and upper**) and prune:
 - If min parent, max children, when $\alpha_{children} \geq \beta_{parent}$, prune!
 - If max parent, min children, when $\alpha_{parent} \geq \beta_{children}$, prune!







Alpha-beta Pruning - Example

MAX


MIN



Alpha-beta Pruning - Example




Alpha-beta Pruning - Practices



● More exercises: http://inst.eecs.berkeley.edu/~cs61b/fa14/ta-materials/apps/ab_tree_practice/

Nondeterministic Game


- In deterministic games with perfect information, Minimax Algorithm gives perfect play
- What if the game is nondeterministic with perfect information?
 - In nondeterministic games, chances are introduced
 - For example:
 - Two people MAX and Min play a game
 - Flip a coin after each player make a decision
 - The result of coin flipping changes the state



Expected Minimax Algorithm

- In nondeterministic games, EXPECTMINIMAX gives perfect play
 - Just like MINIMAX, except we must also handle chance nodes

if *state* is a MAX node then
 return the highest EXPECTIMINIMAX-VALUE of
 SUCCESSORS(*state*)
if *state* is a MIN node then
 return the lowest EXPECTIMINIMAX-VALUE of
 SUCCESSORS(*state*)
if *state* is a chance node then
 return average of EXPECTIMINIMAX-VALUE of
 SUCCESSORS(*state*)




EXPECTIMINIMAX(*s*) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r \underline{P(r)} \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$


Expected Minimax Algorithm

- Unlike MINIMAX algorithm where only **order** of terminal nodes matters, in EXPECTMINIMAX algorithm, **exact values** of terminal nodes also matter!

MINIMAX



EXPECTMINIMAX



Today's Topics

- Constraint Satisfaction Problem (CSP)
 - Formulation of CSPs
 - Backtrack Search
 - Techniques for improving CSP solution
 - Tree-structured CSPs
- Game Playing
 - Formulation as Search
 - Minimax Algorithm
 - Alpha-beta Pruning
 - Expected Minimax for Nondeterministic Games
- **Propositional Logic**

Logic

- Logic: knowledge representation language
 - Represent human knowledge as "**sentences**" (a.k.a *axiom*)
 - **Knowledge base (KB)**: a set of sentences
- Examples
 - **Propositional logic**
 - Boolean logic
 - **First-order logic**
 - Quantifiers \forall, \exists , objects and relations
- Key components in Logic
 - Syntax: how to write sentences
 - Semantics: how to interpret sentences
 - Reasoning/Inference: What new knowledge can be derived from known facts

B or $\neg B$

Propositional Logic - Syntax

Syntax:

- Atomic sentence
 - A single propositional symbol, like A B (A can be True or False)
- Logical connectives
 - \neg not
 - \wedge and (**conjunction**)
 - \vee or (**disjunction**)
 - \Rightarrow (*or* \rightarrow) implication
 - \Leftrightarrow if and only if
- Complex sentence
 - $A \vee B$, $A \vee \neg C \Rightarrow B$, ...
- A special type of sentence: Horn clause

Syntax Forms

Syntax Forms:

CNF (Conjunction Normal Form): $(A \vee \neg B) \wedge (A \vee \neg C \vee D)$ ^{" \wedge "}

- CNF consists of **clauses** that are connected by conjunction.
 - **Clauses:** disjunctions ^{" \vee "} of **literals** (a symbol or its negation).
 $A, \neg A$
- $(A \vee \neg B) \wedge (A \vee \neg C \vee D)$
 - 2 clauses: $(A \vee \neg B), (A \vee \neg C \vee D)$
 - 4 variables: A, B, C, D
 - Literals: $A, \neg B, \neg C, D$

DNF (Disjunction Normal Form): $(A \wedge \neg B) \vee (A \wedge \neg C \wedge D)$
 Δ Δ

- All propositional sentences can be converted to CNF/DNF.
- We will mainly use CNF. For most algorithms, you will need to standardize the sentence by converting it to CNF first.

Horn Clause

Horn Clause:

A Horn clause is a [clause](#) (a [disjunction](#) of [literals](#)) with at most one positive

- $\neg A \vee \neg B \vee \neg C \vee D$
 - $A \wedge B \wedge C \Rightarrow D$

Horn Form: When KB (knowledge base) = **conjunction** of Horn clauses

Why do we care about Horn clause?

- It's a special type! If the sentences are Horn clauses, inference can be done in linear time (exponential for general sentences)

Semantics

- Answers when is a sentence true:

$P \Rightarrow Q$ is equivalent to $\neg P \vee Q$ P \Rightarrow Q And Q \Rightarrow P

P	Q	$\neg P$	$P \wedge Q$	$P \vee Q$	<u>$P \Rightarrow Q$</u>	<u>$P \Leftrightarrow Q$</u>
false	false	true	false	false	true	true
false	true	true	false	true	true	false
true	false	false	false	true	false	false
true	true	false	true	true	true	true

Figure 7.8 Truth tables for the five logical connectives. To use the table to compute, for example, the value of $P \vee Q$ when P is true and Q is false, first look on the left for the row where P is *true* and Q is *false* (the third row). Then look in that row under the $P \vee Q$ column to see the result: *true*.

Questions?

- My slides take the following materials as references:
 - Shirley Chen's slides
 - Yewen Wang's (Winter 2020's TA) slides
 - Prof. Quanquan Gu's (Winter 2020) slides

Thank you!