

Final

First-Order Logic

Objects - numbers, cars, peoples, houses, colors

properties - breezy, tall, red

relations - inside, adjacent, larger, sibling [holds between two objects]

function - father-of, best friend, age [maps object to object]

term - constant or variable or function {Jack, age(Jack), leftof(R)}

Universal Quantification (\forall) - For All

Existential Quantification (\exists) - There Exists

• can switch order of $\forall x \forall y$, $\exists x \exists y$, but NOT for $\forall x \exists y$ $\exists x \forall y$

Uniqueness Quantifier ($\exists!$) - $\exists! x \text{ king}(x)$ - There is only one king x.

Ex: $\forall x \text{ At}(x, \text{UCLA}) \Rightarrow \text{smart}(x)$

$\exists x \text{ At}(x, \text{UCLA}) \Rightarrow \text{tall}(x)$

$\forall y \exists x \text{ Loves}(x, y)$

$\exists x \forall y \text{ Loves}(x, y)$

} different meaning

$\exists x, y \text{ sister}(\text{Spot}, x) \wedge \text{sisiter}(\text{Spot}, y) \wedge \neg(x=y)$

Substitution - exchange x with y in sentence α

$\Rightarrow \forall x \text{ king}(x) \wedge \text{Greedy}(x) \Rightarrow \text{Evil}(x)$

$\text{subst}\{\{x/\text{John}\}, \alpha\}$

$\text{king}(\text{John}) \wedge \text{Greedy}(\text{John}) \Rightarrow \text{Evil}(\text{John})$

• If a sentence is entailed by a FOL knowledge base, then it is entailed by a finite subset of that propositional knowledge base

Unification - replace variables to make two statements equivalent

$\text{Knows}(\text{John}, x) \text{ Knows}(\text{John}, \text{Jane}) \quad \theta = \{x/\text{Jane}\} \rightarrow \text{Knows}(\text{John}, \text{Jane}) \times 2$

Resolution - works the same as in propositional logic

• $\text{KB}(A) \models \alpha(\text{query}) \rightarrow \Delta \wedge \neg \alpha$ (prove a contradiction)

① Convert English sentences to FOL

② Convert FOL to CNF \rightarrow Perform Resolution using Unification as necessary

• If there is a contradiction, then query is proven.

Definite Clause - exactly one positive literal

• $\neg A \vee \neg B \vee C \Rightarrow A \wedge B \Rightarrow C$

Horn Clause - at most one positive literal

• $\neg A \vee \neg B \vee C \Rightarrow A \wedge B \Rightarrow C$

(2)

Conversion to CNF

$$\bullet \exists x \text{ crown}(x)$$

$$\hookrightarrow \text{crown}(c_1)$$

skolem constant

$$\bullet \forall y \exists x \text{ crown}(x)$$

$$\hookrightarrow \forall y \text{ crown}(f(y))$$

skolem function

Replace the \exists variable as
 $F(\text{variable})$

Then remove \exists and \forall terms

① Eliminate $\Rightarrow, \Leftrightarrow$

$$\bullet \alpha \Rightarrow \beta \rightarrow \neg \alpha \vee \beta$$

$$\bullet \alpha \Leftrightarrow \beta \rightarrow (\alpha \Rightarrow \beta) \wedge (\beta \Rightarrow \alpha) \rightarrow (\neg \alpha \vee \beta) \wedge (\neg \beta \vee \alpha)$$

② push \neg inwards

$$\bullet \neg \forall x P \rightarrow \exists x \neg P$$

$$\bullet \neg \exists x P \rightarrow \forall x \neg P$$

De Morgan's

$$\begin{cases} \neg(\alpha \wedge \beta) \rightarrow (\neg \alpha \vee \neg \beta) \\ \neg(\alpha \vee \beta) \rightarrow (\neg \alpha \wedge \neg \beta) \end{cases}$$

③ Skolemize (see above) and drop \forall, \exists terms④ Distribute \vee over \wedge

$$\bullet (\alpha \wedge \beta) \vee \gamma \rightarrow (\alpha \vee \gamma) \wedge (\beta \vee \gamma)$$

Probability

Classical logic is "monotonic" - things that are true will remain true even with the addition of new information (ignores possible contradictions)

Degrees of Belief - how much do I believe that something is true? $[0, 1]$

Rules: ① $0 \leq \text{Pr}(\alpha) \leq 1$

$$\text{② } \alpha \text{ is inconsistent} \Leftrightarrow \text{Pr}(\alpha) = 0$$

$$\text{③ } \alpha \text{ is valid} \Leftrightarrow \text{Pr}(\alpha) = 1$$

$$\text{④ } \text{Pr}(\alpha) + \text{Pr}(\neg \alpha) = 1$$

$$\text{⑤ } \text{Pr}(\alpha \vee \beta) = \text{Pr}(\alpha) + \text{Pr}(\beta) - \text{Pr}(\alpha \wedge \beta) \quad [\alpha, \beta \text{ not mutually exclusive}]$$

$$\text{⑥ } \text{Pr}(\alpha \vee \beta) = \text{Pr}(\alpha) + \text{Pr}(\beta) \quad [\alpha, \beta \text{ are mutually exclusive/independent}]$$

Belief Change - introduction of new information, must re-evaluate probabilities

$$\text{Pr}(w | B) = \begin{cases} 0 & \text{if } w \models \neg B \text{ [if } B \text{ is 0 in that world, change probability to 0]} \\ \frac{\text{Pr}(w)}{\text{Pr}(B)} & \text{if } w \models B \text{ [original probability of world } \div \text{ new total probability after zeroing out stuff]} \end{cases}$$

Bayes Conditioning: $\text{Pr}(\alpha | \beta) = \frac{\text{Pr}(\alpha \wedge \beta)}{\text{Pr}(\beta)}$ [if α, β independent - $\text{Pr}(\alpha \wedge \beta) = \text{Pr}(\alpha) \cdot \text{Pr}(\beta)$]

Independence: $[\text{Pr}(\alpha | \beta) = \text{Pr}(\alpha)]$, $[\frac{\text{Pr}(\alpha \wedge \beta)}{\text{Pr}(\beta)} = \text{Pr}(\alpha)]$ check if $\{P(E) = P(E|B) \& P(B) = P(B|E)\}$


Chain Rule: $\text{Pr}(\alpha_1 \wedge \alpha_2 \wedge \dots \wedge \alpha_n) = \text{Pr}(\alpha_1 | \alpha_2 \dots \alpha_n) \text{Pr}(\alpha_2 | \alpha_3 \dots \alpha_n) \dots \text{Pr}(\alpha_n)$

\hookrightarrow For Graph: use each nodes' parents/children $\rightarrow \text{Pr}(a) \text{Pr}(b) \text{Pr}(c | a) \text{Pr}(d | a, b) \dots$

Case Analysis: Find probability when not explicitly given $\rightarrow P(T) = P(T|0)P(0) + P(T|1)P(1)$ use all possible pairings for needed variables

Bayes Rule: $\text{Pr}(\alpha | \beta) = \frac{\text{Pr}(\alpha \wedge \beta)}{\text{Pr}(\beta)} = \frac{\text{Pr}(\beta | \alpha) \text{Pr}(\alpha)}{\text{Pr}(\beta)}$ $[\text{Pr}(\alpha \wedge \beta) = \text{Pr}(\alpha | \beta) \text{Pr}(\beta)]$

Bayesian Networks



Parents: $A: E, B$ $C: A$ $R: E$ $B: \emptyset$

Descendants: $E: R, A, C$ $C: \emptyset$ $B: A, C$

Non-Descendants: $A: R$ $C: R, E, B$ $E: B$

→ exclude V and its direct parents

Markovian Assumptions: $I(V, \text{Parents}(V), \text{Non-Descendants}(V))$

independence node → write assumptions for all nodes in a graph

→ Every node in a Bayesian Network is conditionally independent of its non-descendants, given its parents. → $I(A, \emptyset, BE)$ means A independent of B, E

d-separation: if any path from $X \rightarrow Y$ is open, then it is not d-separated.

if at least one path from $X \rightarrow Y$ is blocked, then it is d-separated.

dsep^{short known end}(B, EC, R) [known = Z]

① sequential: $\rightarrow w \rightarrow$ [blocked iff $w \in Z$]

② divergent: $\leftarrow w \rightarrow$ [blocked iff $w \in Z$]

③ convergent: $\rightarrow w \leftarrow$ [blocked iff $w \neq Z$ AND descendants(w) $\notin Z$]

Prior Marginal - CPT before we have any evidence (probability a random person has disease)

Posterior Marginal - CPT after acquiring evidence (probability person has disease after test res)

Most Probable Explanation (MPE) - assume the final consequence and search for the most probable query [uses all variables]

Maximum a Posteriori (MAP) - same idea as MAP, but don't include all variables [generally more complex, less efficient]

Same Decision Probability (SDP) - what is the probability that our decision [to prescribe medicine] will stay the same after getting more information [tests]

Complexity of Inference (PIR) - variable elimination, conditioning

• Time/Space complexity: $O(nd^w)$ n : #variables d : #of values $[0,1]$ w : tree width

• normal tree: $w=1$; polytree: nodes have multiple parents, $w = \max$ # of parents a node

Weighted Model Counting (WMC) - more general version of #SAT

• Instead of finding the number of SAT worlds, find the sum of the weights of SAT worlds

• weight of a world is the product of weights assigned to its literals

• compile Δ formula into smooth, decomposable, deterministic MNF circuit → WMC in linear time

Modeling Logic as a Bayesian Network

① Decide variables/values ② Decide edges ③ Create CPTs {cdd, flu, ...} example

• Bipartite - edges only go from the causes to the effects, 2 layers

• complete data - no ? in the data, everything filled out {model with Bayesian Net}

• incomplete data - have at least one ? unknowns in the data {model with Expectation Maximization}

• EM - iterative method to find local maximum or MAP estimates of parameters

Maximum likelihood principle - learn the parameters that maximize the likelihood of observing a certain set of data

calculate likelihood: $\theta_1 = \Pr(e_1) \cdot \Pr(e_2) \cdot \dots \cdot \Pr(e_n)$ [Based on Bayesian Net #1]

$\theta_2 = \Pr_2(e_1) \cdot \Pr_2(e_2) \cdot \dots \cdot \Pr_2(e_n)$ [Based on Bayesian Net #2]

Choose the greater one as your parameters/Bayesian Network

MLP with complete data (given a complete dataset)

① Find Empirical Distribution - for all possibilities of the variables, what is the probability of randomly picking that one from the dataset?

② Find the MLP (ex: $\theta_5(h) = \frac{\Pr(s, h)}{\Pr(h)} = \frac{\Pr_0(s, h)}{\Pr_0(h)} = \frac{10/16}{12/16} = \frac{5}{6}$)

MLP with incomplete data (given an incomplete dataset)

① Use Expectation Maximization (EM)

• basically guess random values for unknown variable and choose the assignment that yields the highest probability; use this in the CPT

• repeat as necessary for any and all unknown variables

• once dataset has been filled up, follow steps for complete dataset

Optimizing only using MLP can lead to overfitting [model is too specific to the training data and doesn't generalize to test data]

Supervised Learning [Query-oriented, Labeled Data] - classification, regression

Unsupervised Learning [Model-oriented, Unlabeled Data] - clustering

Loss Function - evaluates how close test results are to the expected result (i.e. mean-squared)

Arithmetic Circuits - given a query, perform WMC on equivalent arithmetic circuit

• $A = T \{ \lambda_a = 1, \lambda_{\bar{a}} = 0 \}$ $A = F \{ \lambda_a = 0, \lambda_{\bar{a}} = 1 \}$ $A = ? \{ \lambda_a = 1, \lambda_{\bar{a}} = 1 \}$

• Evaluating AC is $O(n)$, but converting query to AC is $O(nd^m)$

• values for the θ stuff in the AC should be given in problem

Cross Entropy (Loss Function): $CE = \sum Q(x) \log_2 P(x)$ $P(x)$: prediction $Q(x)$: label [prefer smaller CE]

Entropy - uncertainty in a single distribution, used for decision trees

Cross Entropy - measure similarity between two distributions, loss function for gradient descent

Machine Learning

Decision Trees and Random Forests - examples of classifiers (supervised)

Entropy - less entropy means more certainty $[ENT(X) = - \sum Pr(x) \log_2 Pr(x)]$

Conditional Entropy - more commonly used

• If we know $ENT(X)$ and $Y=y$, then $ENT(X|Y) = - \sum Pr(x|y) \log_2 Pr(x|y)$

• If we don't know Y but want to observe it, then $ENT(X|Y) = - \sum Pr(y) ENT(X|y)$

Additional information can never increase entropy $\rightarrow ENT(X|Y) \leq ENT(X)$ [information gain]

Decision Tree - interpretable classifier

• best to split into training data and test data to train tree and check accuracy

• Cross-validation: perform the train/test process multiple times but splitting data differently

• determine how to create decision tree / how to split it based on entropy

① Look at all Input Attributes, choose one with lowest conditional entropy based on result/class

[greatest information gain] and split on that

• $ENT(D|A) = Pr(a) \cdot [-Pr(d|a) \log_2 Pr(d|a) + Pr(\bar{d}|a) \log_2 Pr(\bar{d}|a)]$

$ENT(D|\bar{A}) = Pr(\bar{a}) \cdot [-Pr(d|\bar{a}) \log_2 Pr(d|\bar{a}) + Pr(\bar{d}|\bar{a}) \log_2 Pr(\bar{d}|\bar{a})]$

$ENT(D|A) = ENT(D|a) + ENT(D|\bar{a})$ [minimize ENT]

• If CE = 0 when splitting on any variable, immediately split on that one

② Split the examples into True or False for the variable (A) that you split on, based on the class for that example

• For $A = T$: try to split on remaining input attributes

For $A = F$: try to split on remaining input attributes

$[X: A=F, D=YES] \rightarrow A=F, +X,$

③ Keep splitting until leaves of all the tree are either all (+) or all (-)

• Leaves that are all (+) become YES, Leaves that are all (-) become NO

Random Forests - more refined version of decision trees


• Decision Trees are often highly specialized to just one dataset, bad generalization

• Random Forests create multiple copies of training data that include different examples from the data, and create DTs on each copy dataset using randomly selected features (input attributes)

• When new data point enters Random Forest, put it through each little DT and select the final class / result that has a majority vote

Bayesian Network Classifier - set a threshold (T) to classify the inputs

$C = \begin{cases} C & \text{iff } Pr(C|a_1, a_2, \dots, a_N) \geq T \\ \bar{C} & \text{iff } Pr(C|a_1, a_2, \dots, a_N) < T \end{cases}$



X, Y are independent given C

Naive Bayes Classifier - think of spam email example

• these words have probability X of being in normal email, probability Y of being in spam,

open new email and use these probabilities to analyze words \rightarrow decide

⑥

Neural Networks


Neuron - building block of a neural network

Activations - input/outputs of a neuron, typically numbers

Weights (w_{ij}) - associated with input/output activation, value for each input multiplied by weight in calculation

Bias (b) - additional term that may be added in calculation

Activation Function [$g(\cdot)$] - determines if neuron is on/off



$$a_4 = g(2a_1 - a_2 + 0a_3 + b)$$

$\rightarrow a_4$ can be used as input into other neurons

① step function ($t = \text{Threshold}$) ② Sign function ③ sigmoid ④ ReLU

$$\begin{cases} g(x) = 1 & \text{if } x \geq t \\ g(x) = 0 & \text{if } x < t \end{cases}$$

$$\begin{cases} g(x) = 1 & \text{if } x \geq 0 \\ g(x) = -1 & \text{if } x < 0 \end{cases}$$

$$g(x) = \frac{1}{1 + e^{-x}}$$

$$g(x) = \max(0, x)$$

Feed-Forward Neural Network - neuron outputs only feed forward into next neuron


Fully-connected layer - every neuron takes input from all neurons in previous layer

depth - number of layers in neural network (including output layer, excluding input layer)

deep learning - ability to train "deep" (many layers) neural networks

Universal Function Approximators - neural networks can express every function up to a ^{practically expressive} constant error

Neurons can represent basic logic gates as long as they are linearly separable



\rightarrow a full neural network can create more complex, non-linear functions

Training neural networks is like an optimization problem

\rightarrow use loss functions like cross-entropy, mean square error

perform optimization using Gradient Descent - find local min/max of function to

optimize parameters, fit line to data

Accuracy - how well can the neural network classify testing data after

working through the training data?

When data set is too big to analyze all at once!

Epoch - when an entire dataset is passed ^(all the batches) forward and backward

through the neural network only once

Batch - divide the dataset into multiple batches when it's too big

Stopping criteria - split training data into 70% (80%) and validation (20%)

• train on 70% within an epoch, using validation data as test data

• monitor performance on validation data over epochs - when performance peaks/plateaus, we stop there